
Gen AI Hands On- 1

Name: Suchitra Shankar

Class: 6J

SRN: PES2UG23CS608

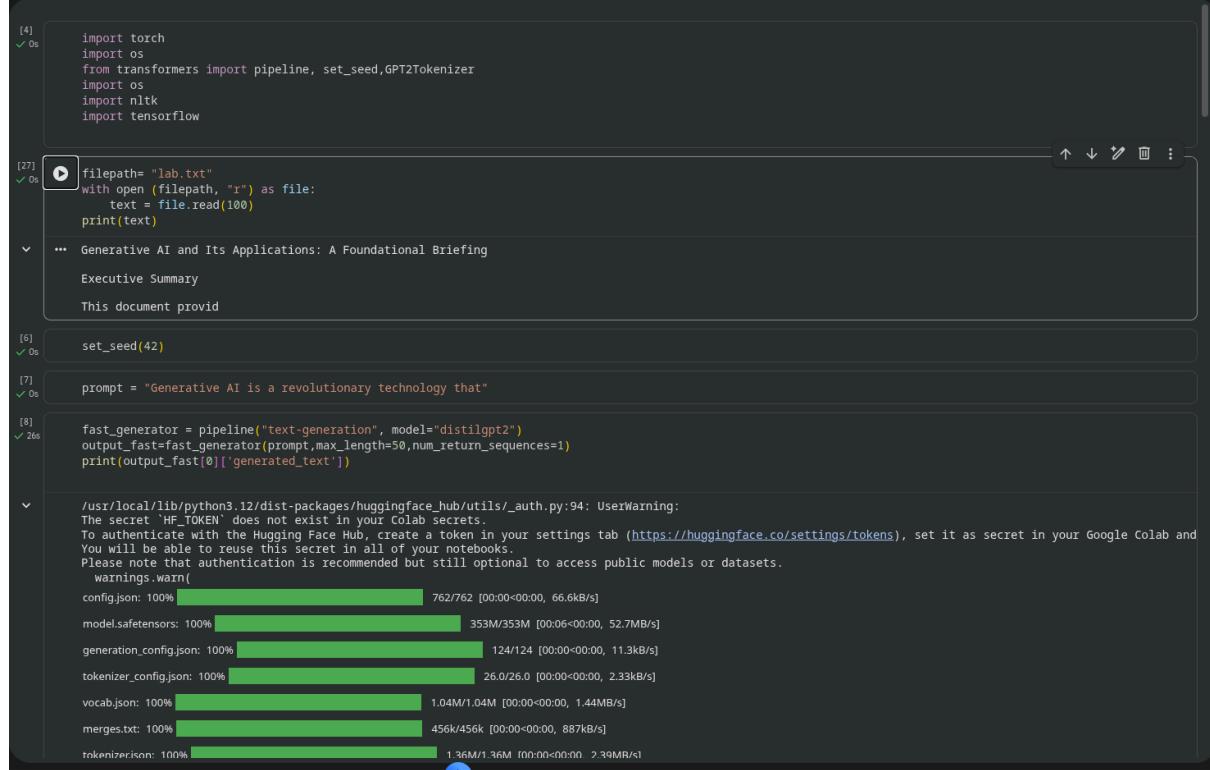
Date: 27 Jan 2026

Email: suchitrashankar07@gmail.com

Phone number: 6366250274

Github link: <https://github.com/SuchitraShankar07/GenAI2026>

Submission 1: Hands on done in class



```
[4]  os      import torch
      import os
      from transformers import pipeline, set_seed,GPT2Tokenizer
      import os
      import nltk
      import tensorflow

[27]  os      filepath= "lab.txt"
      with open (filepath, "r") as file:
          text = file.read(100)
          print(text)

... Generative AI and Its Applications: A Foundational Briefing
Executive Summary
This document provid

[6]  os      set_seed(42)

[7]  os      prompt = "Generative AI is a revolutionary technology that"

[8]  268    fast_generator = pipeline("text-generation", model="distilgpt2")
          output_fast=fast_generator(prompt,max_length=50,num_return_sequences=1)
          print(output_fast[0]['generated_text'])

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret 'HF_TOKEN' does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
config.json: 100% [██████████] 762/762 [00:00<00:00, 66.6kB/s]
model.safetensors: 100% [██████████] 353M/353M [00:06<00:00, 52.7MB/s]
generation_config.json: 100% [██████████] 124/124 [00:00<00:00, 11.3kB/s]
tokenizer_config.json: 100% [██████████] 26.0/26.0 [00:00<00:00, 2.3kB/s]
vocab.json: 100% [██████████] 1.04M/1.04M [00:00<00:00, 1.44MB/s]
merges.txt: 100% [██████████] 456k/456k [00:00<00:00, 887kB/s]
tokenizerison: 100% [██████████] 1.36M/1.36M [00:00<00:00, 2.39MB/s]
```

```

merges.txt: 100% [456k/456k [00:00<00:00, 887kB/s]
tokenizer.json: 100% [1.36M/1.36M [00:00<00:00, 2.39MB/s]

Device set to use cpu
Truncation was not explicitly activated but 'max_length' is provided a specific value, please use 'truncation=True' to explicitly truncate examples to max length.
Setting 'pad_token_id' to 'eos_token_id':50256 for open-end generation.
Both 'max_new_tokens' (=256) and 'max_length' (=50) seem to have been set. 'max_new_tokens' will take precedence. Please refer to the documentation for more information.
Generative AI is a revolutionary technology that is designed to work with existing AI systems. It has been developed by the University of California, Berkeley. It

The research team led by Professor Daniel Kranz, from the University of California, Berkeley, has developed a program to learn how to use the AI to improve the performance of the software. It has been developed by the University of California, Berkeley. The research team developed the program to learn how to use the AI to improve the performance of the software. It has been developed by the University of California, Berkeley.

[9] ✓ lm
fast_generator=pipeline("text-generation",model="gpt2")
output_fast=fast_generator(prompt,max_length=100,num_return_sequences=2)
print(output_fast[0]['generated_text'])

config.json: 100% [665/665 [00:00<00:00, 53.3kB/s]
model.safetensors: 100% [548M/548M [00:06<00:00, 63.6MB/s]
generation_config.json: 100% [124/124 [00:00<00:00, 4.11kB/s]
tokenizer_config.json: 100% [26.0/26.0 [00:00<00:00, 1.09kB/s]
vocab.json: 100% [1.04M/1.04M [00:00<00:00, 1.45MB/s]
merges.txt: 100% [456k/456k [00:00<00:00, 878kB/s]
tokenizer.json: 100% [1.36M/1.36M [00:00<00:00, 1.69MB/s]

Device set to use cpu
Truncation was not explicitly activated but 'max_length' is provided a specific value, please use 'truncation=True' to explicitly truncate examples to max length.
Setting 'pad_token_id' to 'eos_token_id':50256 for open-end generation.
Both 'max_new_tokens' (=256) and 'max_length' (=100) seem to have been set. 'max_new_tokens' will take precedence. Please refer to the documentation for more information.
Generative AI is a revolutionary technology that allows an AI to be used to solve any problem. For example, it is one of the most effective ways to train a human

The goal of the AI is to learn how to respond to an example. By knowing the basic mental models of human behavior, it can then apply those models to complex problems.
The AI has its own set of problems. For example, it can solve problems that would take a human brain and a computer to solve.
This is how intelligent AI, or AI as a whole, can be used in today's tech world.
This technology will allow an AI to learn to respond to a problem, and to learn how to solve it.
This is a great way to build software that has the power to solve problems - to solve problems for a human, or for a machine.
This is a great way to build software that has the power to solve problems - to solve problems for a human, or for a machine.

I want to say that this is a great demonstration of the potential of AI to be used in the real world.

```

Tokenization

```

[10] ✓ ts
tokenizer = GPT2Tokenizer.from_pretrained("gpt2")

[11] ✓ os
sample_sentence="The quick brown fox jumped over the lazy dog."

[12] ✓ os
tokens = tokenizer.tokenize(sample_sentence)
print(f'Tokens:{tokens}')

... Tokens:[‘The’, ‘quick’, ‘brown’, ‘fox’, ‘jumped’, ‘over’, ‘the’, ‘lazy’, ‘dog’, ‘.’]

[13] ✓ os
token_ids = tokenizer.convert_tokens_to_ids(tokens)
print(f'Token IDs: {token_ids}')

Token IDs: [464, 2068, 7586, 21831, 11687, 625, 262, 16931, 3290, 13]

```

POS tagging: part of speech

```

[14] ✓ os
nltk.download('averaged_perceptron_tagger_eng', quiet = True)

True

[15] ✓ os
nltk.download('punkt_tab')

[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]  Unzipping tokenizers/punkt_tab.zip.
True

[16] ✓ os
pos_tags = nltk.pos_tag(nltk.word_tokenize(sample_sentence))
print(f'POS Tags: {pos_tags}')
#
POS Tags: [('The', 'DT'), ('quick', 'JJ'), ('brown', 'NN'), ('fox', 'NN'), ('jumped', 'VBD'), ('over', 'IN'), ('the', 'DT'), ('lazy', 'JJ'), ('dog', 'NN'), ('.', 'PUNKT')]

[17] ✓ 40s
ner_pipeline=pipeline("ner", model = "dbmdz/bert-large-cased-finetuned-conll03-english", aggregation_strategy="simple")
ner_pipeline("My name is Sylvain and I work at Hugging Face")

config.json: 100% [998/998 [00:00<00:00, 38.6kB/s]
model.safetensors: 100% [1.33G/1.33G [00:36<00:00, 32.4MB/s]

```

```

modelfartsensors: 100% 1.55G/1.55G [00:30<00:00, 32.4MB/s]
Some weights of the model checkpoint at dhmdz/bert-large-cased-finetuned-conll03-english were not used when initializing BertForTokenClassification: ['bert.pooler'
- This IS expected if you are initializing BertForTokenClassification from the checkpoint of a model trained on another task or with another architecture (e.g.
- This IS NOT expected if you are initializing BertForTokenClassification from the checkpoint of a model that you expect to be exactly identical (initializing a B
tokenizer_config.json: 100% 60.0/60.0 [00:00<00:00, 5.1kB/s]

vocab.txt: 213k/? [00:00<00:00, 9.96MB/s]
Device set to use cpu
[{'entity_group': 'PER',
 'score': np.float32(0.99846023),
 'word': 'Sylvain',
 'start': 11,
 'end': 18},
 {'entity_group': 'ORG',
 'score': np.float32(0.9917064),
 'word': 'Hugging Face',
 'start': 33,
 'end': 45}]

[18] snippet=text[:1000]
entities = ner_pipeline(snippet)
print(f'{Entity}<25 | {Type}<10 | {Score}<10')
print("-*50")
for entity in entities:
    if entity['score']>.90:
        print(f'{entity["word"]}<25 | {entity["entity_group"]}<10 | {entity["score"]:.2f}'')

Entity | Type | Score
-----|-----|-----
AI | MISC | 0.98
PES University | ORG | 0.99
AI | MISC | 0.98
Large Language Models | MISC | 0.91
LLMs | MISC | 0.90
Transformer | MISC | 0.99

```

Comparative analysis

```

transformer_section = ''
1. Pretraining: This is the initial, computationally expensive phase. The model is trained on a massive, diverse corpus of raw text (billions of tokens from book
2. Finetuning: In this stage, the pre-trained base model is further trained on a smaller, curated dataset specific to a particular task (e.g., medical question a
...
```

```

fast_sum = pipeline("summarization", model="sshleifer/distillbart-cnn-12-6")
res_fast = fast_sum(transformer_section, max_length=100, min_length=30, do_sample=False)
print(res_fast[0]['summary_text'])

config.json: 1.80k/? [00:00<00:00, 118kB/s]
pytorch_model.bin: 100% 1.22G/1.22G [00:49<00:00, 25.7MB/s]
model.safetensors: 100% 1.22G/1.22G [00:50<00:00, 4.40MB/s]
tokenizer_config.json: 100% 26.0/26.0 [00:00<00:00, 222B/s]
vocab.json: 899k/? [00:00<00:00, 8.29MB/s]
merges.txt: 456k/? [00:00<00:00, 3.91MB/s]
Device set to use cpu
The model is trained on a massive, diverse corpus of raw text (billions of tokens from books, websites, etc.) The goal is to develop a broad, general understandin
```

```

smart_sum = pipeline("summarization", model="facebook/bart-large-cnn")
res_smart = smart_sum(transformer_section, max_length=60, min_length=30, do_sample=False)
print(res_smart[0]['summary_text'])

config.json: 1.58k/? [00:00<00:00, 44.1kB/s]
model.safetensors: 100% 1.63G/1.63G [01:39<00:00, 23.5MB/s]
generation_config.json: 100% 363/363 [00:00<00:00, 25.0kB/s]
vocab.json: 899k/? [00:00<00:00, 30.9MB/s]
merges.txt: 456k/? [00:00<00:00, 17.4MB/s]
tokenizer.json: 1.36M/? [00:00<00:00, 35.3MB/s]
Device set to use cpu
The model is trained on a massive, diverse corpus of raw text (billions of tokens from books, websites, etc.) The goal is to develop a broad, general understandin
```

```

qa_pipeline = pipeline("question-answering", model="distilbert-base-cased-distilled-squad")

config.json: 100% 473/473 [00:00<00:00, 34.4kB/s]
model.safetensors: 100% 261M/261M [00:10<00:00, 28.4MB/s]
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 3.05kB/s]
vocab.txt: 100% 213k/213k [00:00<00:00, 1.16MB/s]
tokenizer.json: 100% 436k/436k [00:00<00:00, 1.32MB/s]
Device set to use cpu

```

```

Device set to use cpu

[23]
questions = [
    "What is the fundamental innovation of the Transformer?",
    "What are the risks of using Generative AI?"
]

for q in questions:
    res = qa_pipeline(question=q, context=text[:5000])
    print(f"\nQ: {q}")
    print(f"A: {res['answer']}")

Q: What is the fundamental innovation of the Transformer?
A: to identify hidden patterns, structures, and relationships within the data

Q: What are the risks of using Generative AI?
A: data privacy, intellectual property, and academic integrity

[24]
mask_filler = pipeline("fill-mask", model="bert-base-uncased")

configjson: 100% [██████████] 570/570 [00:00<00:00, 51.0KB/s]
modelsafetensors: 100% [██████████] 440M/440M [00:11<00:00, 28.2MB/s]
Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForMaskedLM: ['bert.pooler.dense.bias', 'bert.pooler.dense.weight',
- This IS expected if you are initializing BertForMaskedLM from the checkpoint of a model trained on another task or with another architecture (e.g. initializing
- This IS NOT expected if you are initializing BertForMaskedLM from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSeq
tokenizer_config.json: 100% [██████████] 48.0/48.0 [00:00<00:00, 3.98kB/s]
vocab.txt: 100% [██████████] 232k/232k [00:00<00:00, 3.31MB/s]
tokenizer.json: 100% [██████████] 466k/466k [00:00<00:00, 1.29MB/s]
Device set to use cpu

[25]
masked_sentence = "The goal of Generative AI is to create new [MASK]."
preds = mask_filler(masked_sentence)

for p in preds:
    print(f"{p['token_str']}: {p['score']:.2f}")

applications: 0.06
ideas: 0.05
problems: 0.05
systems: 0.04
information: 0.03

```

Submission 2: Assignment: comparative analysis of BERT, BART, RoBERTa (filled in the ipynb file and copy pasted from there)

```

[9]
from transformers import pipeline, set_seed
set_seed(42)

[10]
models = {
    "BERT": "bert-base-uncased",
    "RoBERTa": "roberta-base",
    "BART": "facebook/bart-base"
}

Experiment 1: Text Generation

[11]
prompt = "The future of Artificial Intelligence is"

[12]
for name, model_id in models.items():
    print(f"\n{name} ===")
    try:
        generator = pipeline("text-generation", model=model_id)
        out = generator(prompt, max_new_tokens=30)
        print(out[0]["generated_text"])
    except Exception as e:
        print("ERROR:", e)

    ...
    === BERT ===
    If you want to use 'BertLMHeadModel' as a standalone, add 'is_decoder=True.'
    Device set to use cpu
    The future of Artificial Intelligence is.....
    === RoBERTa ===
    If you want to use 'RobertaLMHeadModel' as a standalone, add 'is_decoder=True.'
    Device set to use cpu
    The future of Artificial Intelligence is
    === BART ===
    Some weights of BartForCausalLM were not initialized from the model checkpoint at facebook/bart-base and are newly initialized: ['lm_head.weight', 'model.decoder.
    You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
    Device set to use cpu
    The future of Artificial Intelligence is0303 apart03 humankind03 housing stimuli1110303 discriminatory discriminatoryEstreamFrame discriminatory03

```

~ Masked Language modeling

```

7] 0s
    masked_sentence = "The goal of Generative AI is to [MASK] new content."
    print(f"\n === BERT ===")
    try:
        masker = pipeline("fill-mask", model="bert-base-uncased")
        preds = masker(masked_sentence)
        for p in preds[:5]:
            print(f'{p["token_str"]}: ({p["score"]:.3f})')
    except Exception as e:
        print("ERROR:", e)

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForMaskedLM: ['bert.pooler.dense.bias', 'bert.pooler.dense.weight',
- This IS expected if you are initializing BertForMaskedLM from the checkpoint of a model trained on another task or with another architecture (e.g. initializing
- This IS NOT expected if you are initializing BertForMaskedLM from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSeque

```

```

8] 0s
    === BERT ===
    Device set to use cpu
    create (0.540)
    generate (0.156)
    produce (0.054)
    develop (0.845)
    add (0.018)

9] 2s
    masked_sentence = "The goal of Generative AI is to <mask> new content."
    print(f"\n === RoBERTa ===")
    try:
        masker = pipeline("fill-mask", model="roberta-base")
        preds = masker(masked_sentence)
        for p in preds[:5]:
            print(f'{p["token_str"]}: ({p["score"]:.3f})')
    except Exception as e:
        print("ERROR:", e)

    === RoBERTa ===
    Device set to use cpu
    generate (0.371)
    create (0.368)
    discover (0.084)
    find (0.021)
    provide (0.017)

10] 0s
    ● masked_sentence = "The goal of Generative AI is to <mask> new content."
    print(f"\n === BART ===")
    try:
        masker = pipeline("fill-mask", model="facebook/bart-base")
        preds = masker(masked_sentence)
        for p in preds[:5]:
            print(f'{p["token_str"]}: ({p["score"]:.3f})')
    except Exception as e:
        print("ERROR:", e)

    ...
    === BART ===
    Device set to use cpu
    create (0.078)
    help (0.066)
    provide (0.061)
    enable (0.036)
    improve (0.033)

```

~ Question Answering

```

11] 0s
    ● context = "Generative AI poses significant risks such as hallucinations, bias, and deepfakes."
    question = "What are the risks?"

```

✓ Question Answering

```
[15]: 0s
    context = "Generative AI poses significant risks such as hallucinations, bias, and deepfakes."
    question = "What are the risks?"
```

```
[16]: 13s
    for name, model_id in models.items():
        print(f"\n==== ({name}) ====")
        try:
            qa = pipeline("question-answering", model=model_id)
            res = qa(question=question, context=context)
            print("Answer:", res["answer"])
            print("Score:", round(res["score"], 3))
        except Exception as e:
            print("ERROR:", e)
```

==== BERT ====
Some weights of BertForQuestionAnswering were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight'.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Answer: AI poses significant risks such as hallucinations, bias, and deepfakes
Score: 0.007

==== RoBERTa ====
Some weights of RobertaForQuestionAnswering were not initialized from the model checkpoint at roberta-base and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight'.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Answer: Generative
Score: 0.011

==== BART ====
Some weights of BartForQuestionAnswering were not initialized from the model checkpoint at facebook/bart-base and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight'.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Answer: poses significant risks such as
Score: 0.054

Task	Model	Classification (Success / Failure)	Observation (What actually happened?)	Why did this happen? (Architectural Reason)
Generation	BERT	Failure	Generated placeholder dots instead of meaningful text. Warning indicated it is not configured as a decoder.	BERT is an encoder-only model trained with Masked Language Modeling, not autoregressive next-token prediction.
	RoBERTa	Failure	Returned the prompt unchanged with no continuation.	RoBERTa is also encoder-only and lacks a decoding mechanism for text generation.

Device set to use cpu
Answer: Generative
Score: 0.011

==== BART ====
Some weights of BartForQuestionAnswering were not initialized from the model checkpoint at facebook/bart-base and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight'.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Answer: poses significant risks such as
Score: 0.054

Task	Model	Classification (Success / Failure)	Observation (What actually happened?)	Why did this happen? (Architectural Reason)
Iteration	BERT	Failure	Generated placeholder dots instead of meaningful text. Warning indicated it is not configured as a decoder.	BERT is an encoder-only model trained with Masked Language Modeling, not autoregressive next-token prediction.
	RoBERTa	Failure	Returned the prompt unchanged with no continuation.	RoBERTa is also encoder-only and lacks a decoding mechanism for text generation.
I-Mask	BART	Failure	Generated highly incoherent and noisy text with random tokens and repetitions.	BART is encoder-decoder, but using it via a causal LM pipeline without fine-tuning leads to unstable generation.
	RoBERTa	Success	Correctly predicted words like "create", "generate", and "produce" with high confidence.	RoBERTa is explicitly trained using Masked Language Modeling, making it well-suited for this task.
	BART	Failure	Threw an error due to missing <mask> token.	RoBERTa expects <mask> instead of [MASK]; tokenizer mismatch caused failure.
Question Answering	BERT	Partial Success	Threw an error due to missing <mask> token.	BART also expects <mask> and is not primarily designed for MLM-style prediction.
	RoBERTa	Failure	Returned the full sentence containing the answer, but with a very low confidence score.	Base BERT is not fine-tuned for QA; span prediction weights are randomly initialized.
	BART	Partial Success	Returned an incomplete and irrelevant fragment ("Generative").	Not fine-tuned for QA; encoder representations alone are insufficient for span extraction.
			Returned a plausible answer fragment but incomplete.	Encoder-decoder model forced into extractive QA without QA fine-tuning.

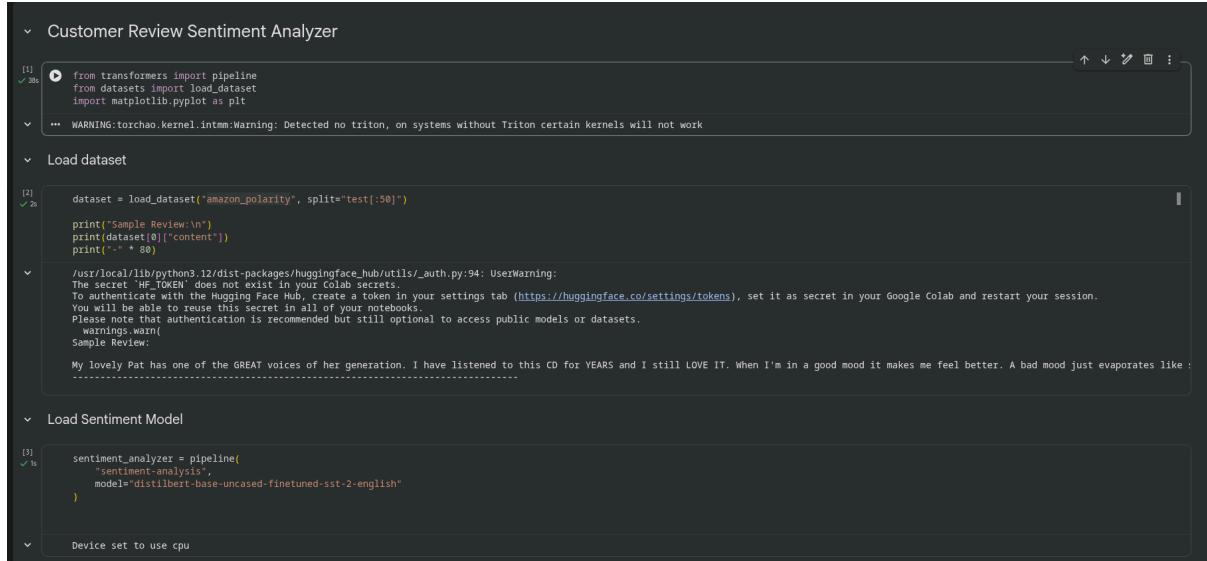
ble-click (or enter) to edit

Task	Model	Classification (Success / Failure)	Observation (What actually happened?)	Why did this happen? (Architectural Reason)
Generation	BERT	Failure	Generated placeholder dots instead of meaningful text. Warning indicated it is not configured as a decoder.	BERT is an encoder-only model trained with Masked Language Modeling, not autoregressive next-token prediction.
	RoBERTa	Failure	Returned the prompt unchanged with no continuation.	RoBERTa is also encoder-only and lacks a decoding mechanism for text generation.
	BART	Failure	Generated highly incoherent and noisy text with random tokens and repetitions.	BART is encoder-decoder, but using it via a causal LM pipeline without fine-tuning leads to unstable generation.
Fill-Mask	BERT	Success	Predicted correct verbs like create and generate with high confidence.	BERT is trained using masked language modeling, so it is naturally good at predicting missing words.

	RoBERTa	Success	Gave meaningful predictions such as generate and create with similar scores.	RoBERTa is similar to BERT and is also trained with masked language modeling, but it requires the <mask> token.
	BART	Success	Predicted reasonable words like create and help, but with lower confidence compared to BERT and RoBERTa.	BART is mainly trained for sequence-to-sequence tasks, so fill-mask works but is not its main strength.
Question Answering	BERT	Partial Success	Returned the full sentence containing the answer, but with a very low confidence score.	Base BERT is not fine-tuned for QA; span prediction weights are randomly initialized.
	RoBERTa	Failure	Returned an incomplete and irrelevant fragment ("Generative").	Not fine-tuned for QA; encoder representations alone are insufficient for span extraction.
	BART	Partial Success	Returned a plausible answer fragment but incomplete.	Encoder-decoder model forced into extractive QA without QA fine-tuning.

PROJECT

Customer Sentiment Analysis:



The screenshot shows a Jupyter Notebook interface with three code cells. The first cell, titled 'Customer Review Sentiment Analyzer', imports necessary libraries (transformers, pipeline, datasets, matplotlib) and loads a dataset ('amazon_polarity'). It prints a sample review and its content. A warning message from the Hugging Face Hub is displayed, stating that the 'HF_TOKEN' secret does not exist in the Colab secrets and providing instructions to create one. The second cell, titled 'Load dataset', shows the dataset loading and printing a sample review. The third cell, titled 'Load Sentiment Model', initializes a sentiment analyzer pipeline using the 'distilbert-base-uncased-finetuned-sst-2-english' model. The output of the notebook shows the sample review and the sentiment analysis result.

```
[1]: 
from transformers import pipeline
from datasets import load_dataset
import matplotlib.pyplot as plt

... WARNING:torchao.kernel.intmm:Warning: Detected no triton, on systems without Triton certain kernels will not work

[2]: 
dataset = load_dataset("amazon_polarity", split="test[:50]")
print("Sample Review:\n")
print(dataset[0]['content'])
print("-" * 88)

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret "HF_TOKEN" does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
Sample Review:

My lovely Pat has one of the GREAT voices of her generation. I have listened to this CD for YEARS and I still LOVE IT. When I'm in a good mood it makes me feel better. A bad mood just evaporates like :
```

```
[3]: 
sentiment_analyzer = pipeline(
    "sentiment-analysis",
    model="distilbert-base-uncased-finetuned-sst-2-english"
)

Device set to use cpu
```

```

✓ Run Sentiment Analysis
[4] 06
texts = [
    str(text)
    for text in dataset["content"]
    if isinstance(text, str) and text.strip() != ""
]

[5] 11b
results = sentiment_analyzer(
    texts,
    truncation=True,
    batch_size=8
)

✓ Predict!
[6] 04
positive = 0
negative = 0

print("\nSentiment Results:\n")

for i, result in enumerate(results):
    label = result["label"]
    score = result["score"]

    if label == "POSITIVE":
        positive += 1
    else:
        negative += 1

    print(f"\tReview {i+1}: ({label}) ({score:.2f})")

...
Sentiment Results:
Review 1: NEGATIVE (0.55)
Review 2: POSITIVE (1.00)
Review 3: NEGATIVE (0.50)
Review 4: NEGATIVE (0.98)
Review 5: POSITIVE (0.97)
Review 6: NEGATIVE (1.00)
Review 7: NEGATIVE (1.00)
...
Review 32: POSITIVE (0.86)
Review 33: NEGATIVE (0.50)
Review 34: POSITIVE (1.00)
Review 35: POSITIVE (1.00)
Review 36: NEGATIVE (1.00)
Review 37: POSITIVE (1.00)
Review 38: NEGATIVE (1.00)
Review 39: NEGATIVE (0.99)
Review 40: POSITIVE (0.99)
Review 41: POSITIVE (0.55)
Review 42: NEGATIVE (1.00)
Review 43: POSITIVE (0.98)
Review 44: POSITIVE (1.00)
Review 45: NEGATIVE (0.99)
Review 46: POSITIVE (0.85)
Review 47: NEGATIVE (1.00)
Review 48: NEGATIVE (1.00)
Review 49: NEGATIVE (1.00)
Review 50: POSITIVE (1.00)

✓ Summary Statistics
[7] 05
print("\n" + "=" * 50)
print("OVERALL SENTIMENT SUMMARY")
print("=" * 50)

print(f"Positive Reviews: {positive}")
print(f"Negative Reviews: {negative}")

=====
OVERALL SENTIMENT SUMMARY
=====
Positive Reviews: 27
Negative Reviews: 23

✓ Visualization
[8] 05
labels = ["Positive", "Negative"]
counts = [positive, negative]

plt.bar(labels, counts)
plt.title("Customer Review Sentiment Distribution")
plt.xlabel("Sentiment")
plt.ylabel("Number of Reviews")
plt.show()

```

Visualization:

Customer Review Sentiment Distribution

